

# Рубежный контроль №2

---

Вариант №4

Студент: Кучеренко М.А.

Группа: ИУ5-21М

## Классификатор №1

LogisticRegression

## Классификатор №2

Multinomial Naive Bayes (MNB)

Необходимо решить задачу классификации текстов на основе любого выбранного Вами датасета (кроме примера, который рассматривался в лекции). Классификация может быть бинарной или многоклассовой. Целевой признак из выбранного Вами датасета может иметь любой физический смысл, примером является задача анализа тональности текста. Необходимо сформировать два варианта векторизации признаков - на основе `CountVectorizer` и на основе `TfidfVectorizer`.

Для каждого метода необходимо оценить качество классификации. Сделайте вывод о том, какой вариант векторизации признаков в паре с каким классификатором показал лучшее качество.

Набор данных - [20 newsgroups text dataset](#)

Классы: 20

Выборка: 18846

```
1 import numpy as np
2 import pandas as pd
3 from typing import Dict, Tuple
4 from sklearn.feature_extraction.text import CountVectorizer,
  TfidfVectorizer
5 from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
6 from sklearn.metrics import accuracy_score, balanced_accuracy_score
7 from sklearn.metrics import precision_score, recall_score, f1_score,
  classification_report
8 from sklearn.metrics import confusion_matrix
9 from sklearn.model_selection import cross_val_score
10 from sklearn.pipeline import Pipeline
11 from sklearn.metrics import mean_absolute_error, mean_squared_error,
  mean_squared_log_error, median_absolute_error, r2_score
12 from sklearn.metrics import roc_curve, roc_auc_score
13 from sklearn.naive_bayes import MultinomialNB
14 from sklearn.linear_model import LogisticRegression
```

```

15 from collections import Counter
16 from sklearn.datasets import fetch_20newsgroups
17 import matplotlib
18 import matplotlib.pyplot as plt
19 import seaborn as sns
20 %matplotlib inline

```

```

1 categories = ["sci.crypt", "sci.electronics", "talk.religion.misc"]
2 newsgroups = fetch_20newsgroups(subset='train', categories=categories)
3 data = newsgroups['data']

```

```

1 def accuracy_score_for_classes(y_true, y_pred):
2     df = pd.DataFrame(data={'t': y_true, 'p': y_pred})
3     res = dict()
4     for c in np.unique(y_true):
5         temp_data_flt = df[df['t'] == c]
6         temp_acc = accuracy_score(
7             temp_data_flt['t'].values,
8             temp_data_flt['p'].values
9         )
10        res[c] = temp_acc
11    return res
12
13 def print_accuracy_score_for_classes(y_true, y_pred):
14     accs = accuracy_score_for_classes(y_true, y_pred)
15     if len(accs) > 0:
16         print('Метка \t Accuracy')
17     for i in accs:
18         print('{} \t {}'.format(i, accs[i]))

```

```

1 # С помощью CountVectorizer преобразуем коллекцию текстовых данных в
   матрицу счётчиков токенов
2 vocabVect = CountVectorizer()
3 vocabVect.fit(data)
4 corpusVocab = vocabVect.vocabulary_
5 print(f'Feature count - {len(corpusVocab)}')

```

```

1 Feature count - 28701

```

```

1 first_el = 0
2 last_el = 9
3
4 for word in list(corpusVocab)[first_el:last_el]:
5     print(f'{word:10}: {corpusVocab[word]}')

```

```
1 from      : 12431
2 mhalld    : 17640
3 lynx      : 16837
4 dac       : 8923
5 northeastern: 18926
6 edu       : 10573
7 mark      : 17150
8 hald      : 13352
9 subject   : 24936
```

```
1 test_features = vocabVect.transform(data)
2 test_features.todense().shape
```

```
1 (1563, 28701)
```

```
1 test_features.todense()
```

```
1 matrix([[0, 0, 0, ..., 0, 0, 0],
2         [0, 0, 0, ..., 0, 0, 0],
3         [0, 0, 0, ..., 0, 0, 0],
4         ...,
5         [0, 0, 0, ..., 0, 0, 0],
6         [1, 0, 0, ..., 0, 0, 0],
7         [0, 0, 0, ..., 0, 0, 0]])
```

```
1 # Cross-validation classification
2 def VectorizeAndClassify(vectorizers_list, classifiers_list):
3     max_acc = 0
4
5     for v in vectorizers_list:
6         for c in classifiers_list:
7             pipeline1 = Pipeline([('vectorizer', v), ('classifier', c)])
8             score = cross_val_score(
9                 pipeline1,
10                 newsgroups['data'],
11                 newsgroups['target'],
12                 scoring='accuracy',
13                 cv=3
14             ).mean()
15             if score > max_acc:
```

```

16         max_acc = score
17         max_v = v
18         max_c = c
19         print(f'Векторизация:\t {v}\nКлассификатор:\t {c}\nAccuracy:\t
{score}')
20         print('='*80)
21
22         print(f'\nЛучший результат: {max_acc}, {type(max_v).__name__},
{type(max_c).__name__}')

```

```

1 vectorizers_list = (
2     CountVectorizer(vocabulary = corpusVocab),
3     TfidfVectorizer(vocabulary = corpusVocab),
4 )
5 classifiers_list = (
6     LogisticRegression(),
7     MultinomialNB(),
8 )
9
10 VectorizeAndClassify(vectorizers_list, classifiers_list)

```

```

1 Векторизация: CountVectorizer(analyzer='word', binary=False,
decode_error='strict',
2     dtype=<class 'numpy.int64'>, encoding='utf-8',
input='content',
3     lowercase=True, max_df=1.0, max_features=None, min_df=1,
4     ngram_range=(1, 1), preprocessor=None, stop_words=None,
5     strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
6     tokenizer=None,
7     vocabulary={'00': 0, '000': 1, '00000000': 2,
'00000000b...
8         '00000001': 4, '00000001b': 5, '00000010': 6,
9         '00000010b': 7, '00000011': 8, '00000011b': 9,
10        '00000100': 10, '00000100b': 11, '00000101':
11        12,
12        '00000101b': 13, '00000110': 14, '00000110b':
13        15,
14        '00000111': 16, '00000111b': 17, '00001000':
15        18,
16        '00001000b': 19, '00001001': 20, '00001001b':
17        21,
18        '00001010': 22, '00001010b': 23, '00001011':
19        24,
20        '00001011b': 25, '00001100': 26, '00001100b':
21        27,
22        '00001101': 28, '00001101b': 29, ...})
23 Классификатор: LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,

```

```

18         intercept_scaling=1, l1_ratio=None, max_iter=100,
19         multi_class='warn', n_jobs=None, penalty='l2',
20         random_state=None, solver='warn', tol=0.0001,
    verbose=0,
21         warm_start=False)
22 Accuracy:  0.9609748411307085
23 =====
24 Векторизация:  CountVectorizer(analyzer='word', binary=False,
    decode_error='strict',
25         dtype=<class 'numpy.int64'>, encoding='utf-8',
    input='content',
26         lowercase=True, max_df=1.0, max_features=None, min_df=1,
27         ngram_range=(1, 1), preprocessor=None, stop_words=None,
28         strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
29         tokenizer=None,
30         vocabulary={'00': 0, '000': 1, '00000000': 2,
    '00000000b...
31             '00000001': 4, '00000001b': 5, '00000010': 6,
32             '00000010b': 7, '00000011': 8, '00000011b': 9,
33             '00000100': 10, '00000100b': 11, '00000101':
    12,
34             '00000101b': 13, '00000110': 14, '00000110b':
    15,
35             '00000111': 16, '00000111b': 17, '00001000':
    18,
36             '00001000b': 19, '00001001': 20, '00001001b':
    21,
37             '00001010': 22, '00001010b': 23, '00001011':
    24,
38             '00001011b': 25, '00001100': 26, '00001100b':
    27,
39             '00001101': 28, '00001101b': 29, ...})
40 Классификатор:  MultinomialNB(alpha=1.0, class_prior=None,
    fit_prior=True)
41 Accuracy:  0.984007541400486
42 =====

```

```

1 Векторизация:  TfidfVectorizer(analyzer='word', binary=False,
    decode_error='strict',
2         dtype=<class 'numpy.float64'>, encoding='utf-8',
3         input='content', lowercase=True, max_df=1.0,
    max_features=None,
4         min_df=1, ngram_range=(1, 1), norm='l2',
    preprocessor=None,
5         smooth_idf=True, stop_words=None, strip_accents=None,
6         sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',
7         tokenizer=None, use...

```

```

8         '00000001': 4, '00000001b': 5, '00000010': 6,
9         '00000010b': 7, '00000011': 8, '00000011b': 9,
10        '00000100': 10, '00000100b': 11, '00000101':
11        12,
12        '00000101b': 13, '00000110': 14, '00000110b':
13        15,
14        '00000111': 16, '00000111b': 17, '00001000':
15        18,
16        '00001000b': 19, '00001001': 20, '00001001b':
17        21,
18        '00001010': 22, '00001010b': 23, '00001011':
19        24,
20        '00001011b': 25, '00001100': 26, '00001100b':
21        27,
22        '00001101': 28, '00001101b': 29, ...})
23
24 Классификатор: LogisticRegression(C=1.0, class_weight=None, dual=False,
25 fit_intercept=True,
26         intercept_scaling=1, l1_ratio=None, max_iter=100,
27         multi_class='warn', n_jobs=None, penalty='l2',
28         random_state=None, solver='warn', tol=0.0001,
29         verbose=0,
30         warm_start=False)
31
32 Accuracy: 0.9366625786143953
33
34 =====
35 =====
36
37 Векторизация: TfidfVectorizer(analyzer='word', binary=False,
38 decode_error='strict',
39         dtype=<class 'numpy.float64'>, encoding='utf-8',
40         input='content', lowercase=True, max_df=1.0,
41         max_features=None,
42         min_df=1, ngram_range=(1, 1), norm='l2',
43         preprocessor=None,
44         smooth_idf=True, stop_words=None, strip_accents=None,
45         sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',
46         tokenizer=None, use...
47         '00000001': 4, '00000001b': 5, '00000010': 6,
48         '00000010b': 7, '00000011': 8, '00000011b': 9,
49         '00000100': 10, '00000100b': 11, '00000101':
50        12,
51        '00000101b': 13, '00000110': 14, '00000110b':
52        15,
53        '00000111': 16, '00000111b': 17, '00001000':
54        18,
55        '00001000b': 19, '00001001': 20, '00001001b':
56        21,
57        '00001010': 22, '00001010b': 23, '00001011':
58        24,
59        '00001011b': 25, '00001100': 26, '00001100b':
60        27,

```

```
39         '00001101': 28, '00001101b': 29, ...})
40 Классификатор: MultinomialNB(alpha=1.0, class_prior=None,
41   fit_prior=True)
42 Accuracy: 0.8675669621264962
43 =====
44 Лучший результат: 0.984007541400486, CountVectorizer, MultinomialNB
```